

Assignment 3

BIOM 5405 - Winter 2017

Name: Maryam Kaka 100929992

Date Submitted: February 23, 2017

All code used to produced plots and obtain values for this assignment was written in python. Five python libraries (`math`, `pandas`, `numpy`, `sklearn` and `matplotlib`) were also included to aid with different aspects of data manipulation and visualization. All the code used can be found in Appendix A. Code output can be found in Appendix B

1.0 - Classifier Scores

- (i) (a) Tables I show the confusion matrices for a classifier assuming 100 samples in each class

Table I: Confusion matrix of classifier with 100 samples in each class

		Predicted	
		T	F
Actual	T	65	35
	F	38	62

- (b) Tables II show the confusion matrices for a classifier assuming 100 positive samples and 1000 negative classes.

Table II: Confusion matrix of classifier with 100 positive samples and 1000 negative

		Predicted	
		T	F
Actual	T	65	35
	F	445	655

- (c) Tables III show the confusion matrices for a classifier assuming 400 samples in each class

Table III: Confusion matrix of classifier with 400 samples in each class

		Predicted	
		T	F
Actual	T	260	140
	F	152	248

Table IV summarizes the result of a chi squared test performed for each case. It can be noted that for all there cases the p-value obtained was much less than 0.05 and therefore each of the classifiers performed better than random. This can also be seen in the precision-recall plot, Figure I, as each of the points are above the respective random classifier line

Table IV: Summary of chi-squared test performed for each case

Case	χ^2	P-Value
A	14.59	1.33×10^{-4}
B	27.51	1.56×10^{-7}
C	58.37	2.16×10^{-14}

- (ii) Table V summarized the precision or positive predictive value (PPV) calculated for each of the cases. Equations 1 was used to calculate the values for each case.

$$PPV = \frac{TP}{TP + FP} \quad (1)$$

Table V: Summary of PPV calculated for each case

Case	PPV
A	0.631
B	0.146
C	0.631

- (iii) (iv) Figure I shows the precision recall curve plotted for each of the cases as well as the curve for a random classifier for each of the cases. Note that a and c on the plot are overlapping.

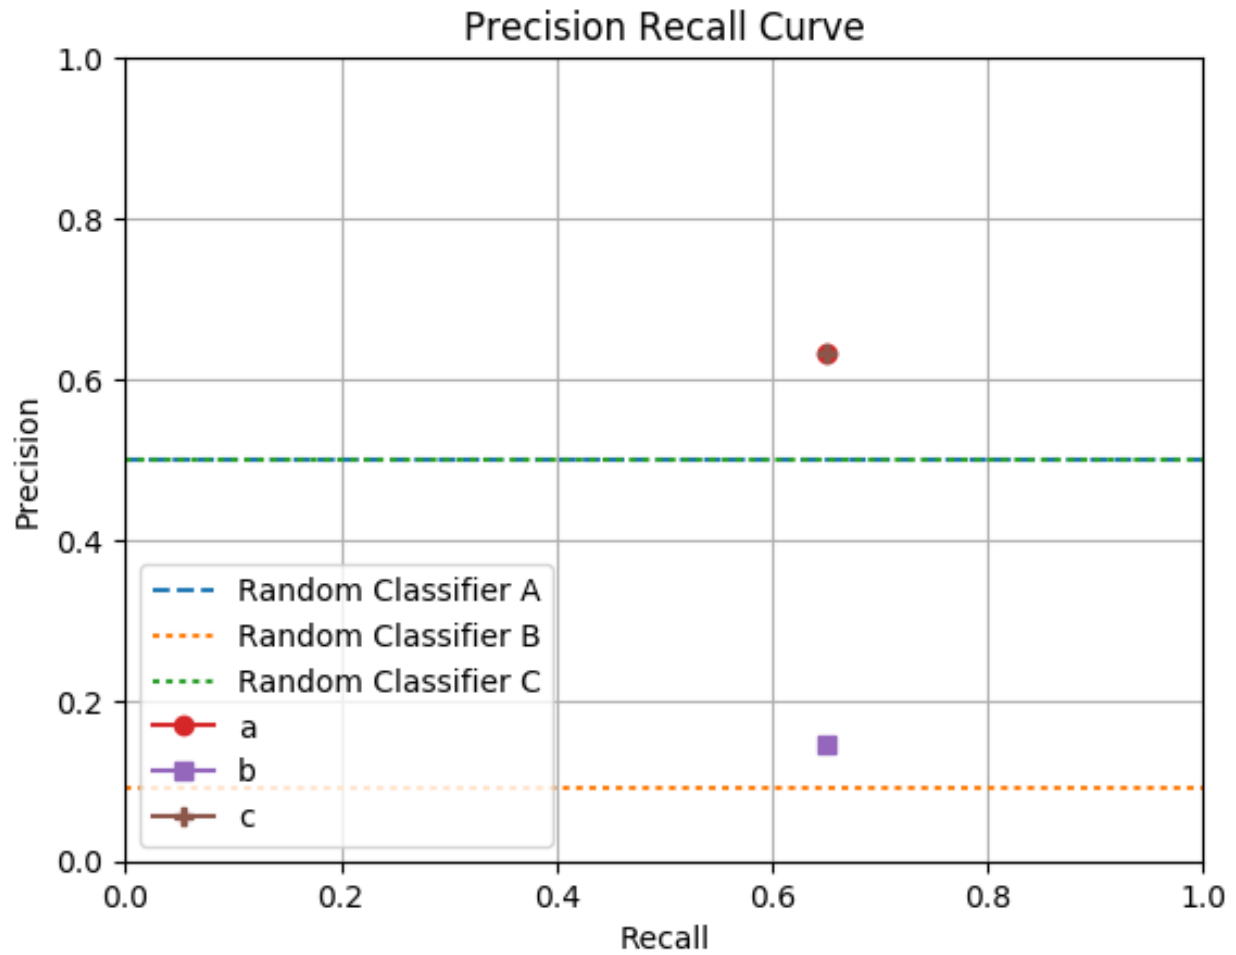


Figure I: Precision Recall curve of each of random classifier performance as well each classifier for each case

2.0

- (i) Figure II shows the receiver operating characteristic curve for the given dataset. The area under the curve was found to be 0.709494.




Figure II: Plot of receiver operating characteristic curve for the collected data

- (ii) Looking at Figure II it can be seen that in order to obtain a sensitivity (i.e. TPR) of at least 75% the maximum specificity (i.e. 1-FPR) that can be obtained is 63%
- (iii) (iv) Figure III shows the precision recall curve obtained. It can be seen on the plot that in order to obtain a sensitivity (i.e. recall) of 75% then a precision of 66.96 % is required.




Figure III: Plot of precision recall curve with 75% sensitivity marked

- (v) Using a bootstrap with 2000 bootstrapped samples resulted in a 95% confidence interval of [0.666, 0.674].
- (vi) A class imbalance of 1000:1 would result in a precision of 0.1% for a recall of 75%. This value was calculated by generating a new dataset (through random sampling), based on the original, which reflects the class imbalance.
- (vii) The fall discovery rate (FDR) is a measure of the number of false positive (i.e. Type I error) over the number of positive predictions

Appendix A - Python Code

```
import pandas as pd
import numpy as np
import scipy.stats as stats
import math
import matplotlib.pyplot as plt
import bisect
from sklearn import metrics
from sklearn import utils

outputLocation = 'Assignment #3/images/'

def chiSquared(data):
    df = ((data.shape[0] - 1) - 1) * ((data.shape[1] - 1) - 1)
    contingencyTable = np.outer(data.RowTotal[0:-1], data.ix["ColTotal"][0:-1])\
        /data.ix["ColTotal", "RowTotal"]
    contingencyTable = pd.DataFrame(contingencyTable, index=rows, columns = col)

    x, p = stats.chisquare(data.ix[0:-1, 0:-1], contingencyTable, ddof = df)
    x = x.sum()
    p = 1 - stats.chi2.cdf(x, df)

    print('Degrees of Freedom: ' + str(df))
    print("Chi^2: " + str(x))
    print("p-value: " + str(p))

def confusionMatrix(tp, fn, fp, tn):
    data = pd.DataFrame([[tp, fn], [fp, tn]], index=rows, columns=col)
    data = data.append(pd.Series(data.sum(), name="ColTotal"))
```

```
data = data.assign(RowTotal=pd.Series(data.sum(axis=1)))
print(data)
return data

def bootstrap(data, nSamples):
    stats = np.empty(nSamples)
    for i in range(0, nSamples):
        sample = utils.resample(data)
        stats[i] = findPrecision(sample['class'], sample['score'])
    np.sort(stats)
    return [stats[math.floor(nSamples*0.025)], stats[math.ceil(nSamples*0.975)]]

def findPrecision(c, s):
    precision, recall, t = metrics.precision_recall_curve(c, s, pos_label=1)
    points = np.array([[i, x] for i,x in enumerate(recall) if x >= 0.75])
    points = points[np.argsort(points[:, 1])]
    return precision[points[0][0]]

print("QUESTION 1")
col = ["TestT", "TestF"]
rows = ["ActualT", "ActualF"]
SENS = 0.65
SPEC = 0.62

ppv = lambda tp, fp: tp/(tp+fp)

print('(i)')
print('a.')
nPos = 100; nNeg = 100
tp = SENS*nPos; fn = nPos - tp; tn = SPEC*nNeg; fp = nNeg - tn
```



```
PPVa = ppv(tp, fp)
chiSquared(confusionMatrix(tp, fn, fp, tn))

print('\nb.')
nPos = 100; nNeg = 1000
tp = SENS*nPos; fn = nPos - tp; tn = SPEC*nNeg; fp = nNeg - tn
PPVb = ppv(tp, fp)
chiSquared(confusionMatrix(tp, fn, fp, tn))

print('\nc.')
nPos = 400; nNeg = 400
tp = SENS*nPos; fn = nPos - tp; tn = SPEC*nNeg; fp = nNeg - tn
PPVc = ppv(tp, fp)
chiSquared(confusionMatrix(tp, fn, fp, tn))

print('\n(ii)')
print('a. PPV: ' + str(PPVa))
print('b. PPV: ' + str(PPVb))
print('c. PPV: ' + str(PPVc))

plt.figure()
plt.plot([0, 1],[0.5, 0.5], linestyle='dashed', label='Random Classifier A')
plt.plot([0, 1],[100/1100, 100/1100], linestyle='dotted', label='Random Classifier B')
plt.plot([0, 1],[0.5, 0.5], linestyle='dotted', label='Random Classifier C')
plt.plot(SENS, PPVa, marker='o', label='a')
plt.plot(SENS, PPVb, marker='s', label='b')
plt.plot(SENS, PPVc, marker='P', label='c')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.title('Precision Recall Curve')
```

```
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.legend(loc = 'lower left')
plt.grid(True)
plt.savefig(outputLocation + 'precision-recall.png', bbox_inches='tight')

print('\nQUESTION 2')
data = pd.read_csv("Assignment #3/assigData3.tsv", sep='\t', index_col=False,
                    header=None, names=["score", "class"])

print('(i)')
fpr, tpr, thresh = metrics.roc_curve(data['class'], data['score'], pos_label=1)
auc = metrics.auc(fpr, tpr)
print('AUC: ' + str(auc))

plt.figure()
plt.plot(fpr, tpr, label='ROC Curve (AUC = %0.2f)' %auc)
plt.title('Reciever Operating Charateristic Curve')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend(loc = 'lower right')
plt.grid(True)
plt.savefig(outputLocation + 'ROC.png', bbox_inches='tight')

print('(ii)')
minSens = 0.75
index = bisect.bisect_left(tpr, minSens)
print('Max Specificity: ' + str(1-fpr[index]))

print('(iii)')
```

```
precision, recall, t = metrics.precision_recall_curve(data['class'],
    data['score'], pos_label=1)
auc = metrics.auc(recall[1:-1], precision[1:-1])

print('(iv)')
index = [i for i,x in enumerate(recall) if x == 0.75][0]
print('Precision (@75'+ '%' + ' Sensitivity): ' + str(findPrecision(data['class'],
    data['score'])))

plt.figure()
plt.plot(recall[1:-1], precision[1:-1],
    label='Precision-Recall Curve (AUC = %0.2f)' %auc)
plt.plot(recall[index], precision[index], marker='o', label='75% Sensitivity Point')
plt.title('Precision Recall Curve')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.legend(loc = 'lower right')
plt.grid(True)
plt.savefig(outputLocation + 'precision-recall2.png', bbox_inches='tight')

print('(v)')
print(bootstrap(data, 2000))

print('(vi)')
modifiedSet = utils.resample(data[data['class'] == 0], n_samples=1000)
modifiedSet = modifiedSet.append(data[data['class'] == 1][:1])
print('Expected Precision: ' + str(findPrecision(modifiedSet['class'],
    modifiedSet['score'])))
```

Appendix B - Code Output

QUESTION 1

(i)

a.

	TestT	TestF	RowTotal
ActualT	65.0	35.0	100.0
ActualF	38.0	62.0	100.0
ColTotal	103.0	97.0	200.0

Degrees of Freedom: 1

Chi²: 14.5931338204

p-value: 0.000133399716113

b.

	TestT	TestF	RowTotal
ActualT	65.0	35.0	100.0
ActualF	380.0	620.0	1000.0
ColTotal	445.0	655.0	1100.0

Degrees of Freedom: 1

Chi²: 27.5117934643

p-value: 1.5613949178e-07

c.

	TestT	TestF	RowTotal
ActualT	260.0	140.0	400.0
ActualF	152.0	248.0	400.0
ColTotal	412.0	388.0	800.0

Degrees of Freedom: 1

Chi²: 58.3725352818

p-value: 2.16493489802e-14

(ii)

a. PPV: 0.6310679611650486

b. PPV: 0.14606741573033707

c. PPV: 0.6310679611650486

QUESTION2

(i)

AUC:0.709494

(ii)

MaxSpecificity:0.63

(iii)

(iv)

Precision(@75%Sensitivity):0.669642857143

(v)

[0.67393278837420523,0.66607773851590102]

(vi)

Expected Precision: 0.00102249488753